

```
In [ ]: !pip3 install torch torchvision torchaudio pandas pyYAML tqdm seaborn opencv-python
```

```
Requirement already satisfied: torch in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (1.12.0)
Requirement already satisfied: torchvision in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (0.13.0)
Requirement already satisfied: torchaudio in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (0.12.0)
Requirement already satisfied: pandas in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (1.4.3)
Requirement already satisfied: pyYAML in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (6.0)
Requirement already satisfied: tqdm in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (4.64.0)
Requirement already satisfied: seaborn in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (0.11.2)
Requirement already satisfied: opencv-python in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (4.6.0.66)
Requirement already satisfied: numpy in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (1.23.1)
Requirement already satisfied: typing-extensions in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from torch) (4.3.0)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from torchvision) (9.2.0)
Requirement already satisfied: requests in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from torchvision) (2.28.1)
Requirement already satisfied: pytz>=2020.1 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from pandas) (2022.1)
Requirement already satisfied: python-dateutil>=2.8.1 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from pandas) (2.8.2)
Requirement already satisfied: colorama in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from tqdm) (0.4.5)
Requirement already satisfied: matplotlib>=2.2 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from seaborn) (3.5.2)
Requirement already satisfied: scipy>=1.0 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from seaborn) (1.8.1)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from matplotlib>=2.2->seaborn) (1.4.4)
Requirement already satisfied: pyparsing>=2.2.1 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from matplotlib>=2.2->seaborn) (3.0.9)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from matplotlib>=2.2->seaborn) (4.34.4)
Requirement already satisfied: cycler>=0.10 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from matplotlib>=2.2->seaborn) (0.11.0)
Requirement already satisfied: packaging>=20.0 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from matplotlib>=2.2->seaborn) (21.3)
Requirement already satisfied: six>=1.5 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from python-dateutil>=2.8.1->pandas) (1.16.0)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from requests->torchvision) (1.26.10)
Requirement already satisfied: certifi>=2017.4.17 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from requests->torchvision) (2022.6.15)
Requirement already satisfied: idna<4,>=2.5 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from requests->torchvision) (3.3)
Requirement already satisfied: charset-normalizer<3,>=2 in c:\users\jmess\appdata\local\programs\python\python39\lib\site-packages (from requests->torchvision) (2.1.0)
```

```
In [ ]: %matplotlib inline
```

```
In [ ]: import torch
import cv2
from matplotlib import pyplot as plt
```

```
In [ ]: model = torch.hub.load('ultralytics/yolov5', 'yolov5s')
```

Using cache found in C:\Users\jmess\.cache\torch\hub\ultralytics\_yolov5\_master  
YOLOv5 2022-7-18 Python-3.9.13 torch-1.12.0+cpu CPU

Fusing layers...

YOLOv5s summary: 213 layers, 7225885 parameters, 0 gradients

Adding AutoShape...

```
In [ ]: def load_image(path):
return cv2.imread(path)[: , : , ::-1]
```

```
In [ ]: def process_image(img):
result = model([img], size=640)

print(result.pandas().xyxy[0])

return result
```

```
In [ ]: IMGS_PATH = ['imgs/all.png', 'imgs/people_1.png', 'imgs/people_2.png']
```

```
In [ ]: IMGS = []

for i in IMGS_PATH:
IMGS.append(load_image(i))
```

a)

```
In [ ]: def draw_result(frame, results):
frame = frame.copy()
for box in results.xyxy[0]:
if int(box[5]) == 0:
xB = int(box[2])
xA = int(box[0])
yB = int(box[3])
yA = int(box[1])

cv2.rectangle(frame, (xA, yA), (xB, yB), (255, 0, 0), 2)

return frame
```

```
In [ ]: def process(img):
plt.imshow(img)
plt.show()

result = process_image(img)

print(result)

plt.imshow(draw_result(img, result))
plt.show()
```

```
In [ ]: process(IMGS[0])
```



	xmin	ymin	xmax	ymax	confidence	class \
0	726.216492	870.891785	967.580078	1088.871460	0.846512	2
1	1104.301514	841.032959	1196.471558	923.366333	0.829974	2
2	393.870117	922.003052	694.894958	1192.261230	0.800172	2
3	1.707250	910.309021	84.191330	1184.690308	0.776526	0
4	636.185120	826.378418	760.100464	929.397766	0.770256	2
5	1045.868408	804.348999	1140.609253	883.489014	0.762862	2
6	1600.351807	493.414764	1672.191406	587.836914	0.680482	9
7	256.838928	852.810669	328.951660	1052.742432	0.631074	0
8	447.830383	725.431458	625.831116	910.149231	0.628749	5
9	1189.572998	697.095337	1907.240234	1186.575684	0.597191	7
10	1206.703735	593.303833	1252.806030	651.825256	0.585567	9
11	878.548706	793.447327	911.642090	870.598999	0.522277	0
12	971.822571	797.682190	1003.063477	869.149353	0.472607	0
13	718.716797	753.658752	812.098267	850.487915	0.451698	2
14	396.205566	838.769409	441.134613	986.398010	0.450832	0
15	507.221191	523.475037	573.016846	602.472778	0.437403	9
16	1020.958191	670.210266	1041.742432	708.367859	0.422663	9
17	392.386963	924.273560	691.774841	1186.112305	0.386347	7
18	1196.590454	700.699097	1843.417358	1188.745605	0.296920	5
19	964.837463	767.659180	1010.002258	861.787720	0.285839	0

	name
0	car
1	car
2	car
3	person
4	car
5	car
6	traffic light
7	person
8	bus
9	truck
10	traffic light
11	person
12	person
13	car
14	person
15	traffic light
16	traffic light
17	truck
18	bus
19	person

image 1/1: 1235x1920 6 persons, 6 cars, 2 buss, 2 trucks, 4 traffic lights

Speed: 14.9ms pre-process, 116.6ms inference, 16.9ms NMS per image at shape (1, 3, 416, 640)



Como podemos ver, já que existem varios elementos na foto boa parte das pessoas foi encontrada pelo modelo. Mas dá para ver que quanto mais profundo na foto o elemento está menor a confiança na classificação temos. Provavelmente devido menor informação de pixels para representá-lo dentro da foto. A rede identificou todas as pessoas na foto sem nenhuma falha.

```
In [ ]: process(IMGs[1])
```



	xmin	ymin	xmax	ymax	confidence	class	\
0	528.242798	213.827530	608.271057	445.049377	0.859260	0	
1	369.474243	211.207169	384.192078	257.956268	0.816155	0	
2	392.125702	212.672897	405.924194	254.231247	0.733099	0	
3	99.137474	249.691849	149.096802	408.160004	0.711861	0	
4	519.139893	276.049927	558.115784	353.744110	0.657766	26	
5	48.602715	255.429077	111.330231	410.260864	0.655491	0	
6	445.193054	199.727753	473.608276	257.162567	0.448671	58	
7	582.135010	246.941681	621.685974	377.284088	0.395870	58	
8	416.761749	210.722229	428.724884	239.429962	0.371667	0	
9	34.038357	305.115082	92.542786	430.863617	0.362988	56	
10	582.520630	325.183167	617.694824	377.309143	0.335677	58	
11	659.666870	109.544678	700.000000	153.703491	0.307373	58	
12	521.921143	209.655411	605.688599	448.508728	0.261612	26	
13	36.373814	298.593567	98.360062	431.164307	0.258987	0	

	name
0	person
1	person
2	person
3	person
4	handbag
5	person
6	potted plant
7	potted plant
8	person
9	chair
10	potted plant
11	potted plant
12	handbag
13	person

image 1/1: 525x700 7 persons, 2 handbags, 1 chair, 4 potted plants

Speed: 4.0ms pre-process, 121.6ms inference, 1.0ms NMS per image at shape (1, 3, 480, 640)



Uma rua com uma quantidade relativa de objetos podemos ver que todas as pessoas foram identificadas, contudo o modelo acabou errado e identificado uma pessoa duas vezes.

In [ ]: process(IMGs[2])



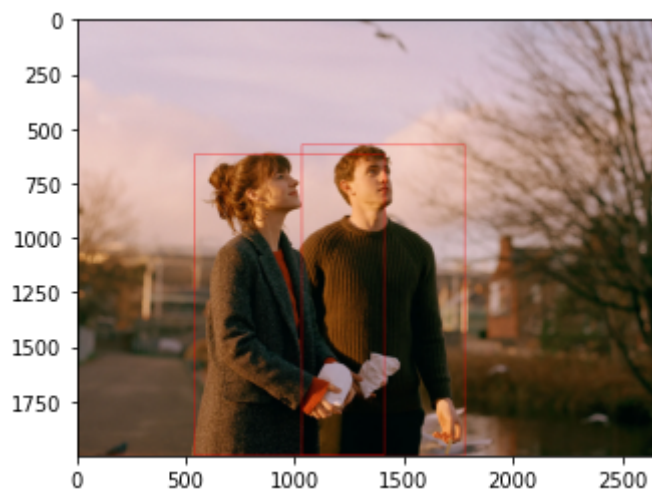


	xmin	ymin	xmax	ymax	confidence	class \
0	1028.121338	574.823120	1774.388672	1998.000000	0.915181	0
1	537.640442	618.084961	1406.829224	1992.961426	0.865246	0

```

name
0 person
1 person
image 1/1: 1998x2664 2 persons
Speed: 24.9ms pre-process, 111.6ms inference, 1.0ms NMS per image at shape (1, 3, 480, 640)

```



Os dois unico elementos da foto propositalmente são duas pessoas e podemos ver que o modelo o identifica, com um nivel consideravel de confiança as duas pessoas.

**b)**

```

In [ ]: # Transforma o video em imagens para ser usado no modelo

def get_frames(path):
    vidcap = cv2.VideoCapture(path)
    vidcap.set(cv2.CAP_PROP_FRAME_WIDTH, 1280)
    vidcap.set(cv2.CAP_PROP_FRAME_HEIGHT, 720)
    lista_results = []

    success, image = vidcap.read()
    while success:
        image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
        lista_results.append(image)

```

```
        success,image = vidcap.read()

    return lista_results
```

```
In [ ]: # Pega todos os frames
results = get_frames('videos/cross.mp4')
```

```
In [ ]: # Executa os frames no modelo

t = model(results)
pd = t.pandas().xyxy
```

```
In [ ]: # Verifica e calcula os centros dos quadrados demarcadores
def get_path(pd):
    path = []

    for j in pd:
        i = j[j['class'] == 0]

        if len(i) > 0:
            i = i.iloc[0]
            x = (i['xmax'] + i['xmin']) // 2
            y = (i['ymax'] + i['ymin']) // 2
            path.append((int(x), int(y)))

    return path
```

```
In [ ]: # Pinta área a partir dos centros dos quadrados demarcadores calculado
def print_area(img, point, w):
    for i in range(point[0] - w, point[0] + w):
        for j in range(point[1] - w, point[1] + w):
            try:
                img[j][i] = (255, 0, 0)
            except IndexError:
                pass
```

```
In [ ]: # Angraria os centros
a = get_path(pd)
```

```
In [ ]: # Pinta o caminho na imagem

img = results[0].copy()
for i in a:
    print_area(img, i, 30)
```

```
In [ ]: # Exibe a imagem

plt.imshow(img)
plt.show()
```

